

3 - Matrices, Optimal Portfolios and Factors

Espen Sirnes

2024-09-18

Table of contents

1	Matrices	2
2	Algebra with Matrices	2
2.1	Matrix Multiplication	3
2.2	Adding and Subtracting Matrices	4
2.3	Dividing with a Matrix	5
2.4	Solving Equations with Matrix Algebra	6
2.5	Transposing	7
3	Calculus and matrices	8
4	Optimal portfolios with more than one asset	9
4.1	Optimal Portfolios with Any Number of Assets	9
5	Empirical example - optimal portfolio and the portfolio front	11
5.1	The historic mean and covariance matrix	13
5.2	Plotting the portfolio front	15
6	Factors	20
6.1	Construction of Factors	20
6.2	Factor Model	21
6.3	Factors in Portfolio Evaluation	21
6.4	Historical Context	21
6.4.1	Coding Challenges:	22
6.4.2	Tentative answer to Challenge 2:	23

This lecture explores the strategic behavior of an investor in the stock market, particularly under the assumption of risk aversion, as discussed in the previous note on utility theory. Risk lovers generally prefer the most risky assets, while risk-neutral investors opt for assets with the highest returns. In contrast, a risk-averse investor seeks to maximize returns without disproportionately increasing volatility, typically measured as variance.

1 Matrices

To calculate optimal portfolios for any number of assets, a basic understanding of matrix algebra is essential. Matrix algebra simplifies the resolution of several equations simultaneously, a process that becomes increasingly complex with the addition of variables. Using matrix functions in software like Excel and various statistical packages allows us to solve systems of equations efficiently without manually computing each one.

Matrices not only streamline the computation but also simplify notation, making the formulation of equations for optimal portfolios more manageable.

A matrix is a structured array of numbers arranged in rows and columns, essentially a set of vectors. Here's an example of a vector:

```
import numpy as np
np.random.randint(0, 100, 3)
```

```
array([21, 58, 41])
```

Combining several vectors side-by-side forms a matrix:

```
np.random.randint(0, 100, (2, 3))
```

```
array([[20,  0, 81],
       [47, 49, 48]])
```

This format is sometimes denoted as $\mathbf{X}_{N \times K}$ to indicate the number of rows (N) and columns (K).

2 Algebra with Matrices

Matrix algebra operates under similar principles to ordinary algebra—allowing addition, subtraction, multiplication, and division (through inversion)—but it also

requires adherence to specific rules.


2.1 Matrix Multiplication

The core operation in matrix algebra is matrix multiplication, which combines elements from the rows of the first matrix with the columns of the second. For example, multiplying a 2×3 matrix by a 3×2 matrix yields:

```
X = np.random.randint(0, 5, (2, 3))
Y = np.random.randint(0, 5, (3, 2))
result = np.dot(X, Y)
print(X)
print(Y)
print(result)
```

```
[[3 4 1]
 [2 4 2]]
[[1 0]
 [3 3]
 [1 4]]
[[16 16]
 [16 20]]
```

What happens is that we sum the product of the elements in each row of the first matrix and each column of the second. You can for example check that element $[0,0]$ of the result is the sum of the product of the first row of the first matrix, and the first column of the second. An easy way to remember this is to think of the multiplication of $A \times B$ is to follow the lines of the letters:



Due to the rules for matrix multiplication, it requires the number of columns in the first matrix to match the number of rows in the second.

The matrix multiplication is different from the normal multiplication in Python. Normal multiplication can be done with the normal multiplication operator `*`. It will then multiply each element in X with the corresponding element of Y, and both matrices must be of the same size:

```
X = np.random.randint(0, 5, (2, 3))
Y = np.random.randint(0, 5, (2, 3))
result = X*Y
print(X)
print(Y)
print(result)
```

```
[[1 1 0]
 [2 4 1]]
[[2 2 0]
 [4 4 4]]
[[ 2  2  0]
 [ 8 16  4]]
```

The reason for using the former method, is that the former is required for solving sets of equations.

2.2 Adding and Subtracting Matrices

Adding or subtracting matrices is straightforward; simply add or subtract corresponding elements. In Python, the multiplication requires numpy function, but if the matrices are numpy variables, subtraction and addition can be done with the normal operators.

```
import numpy as np

X = np.random.randint(0, 100, (2, 2))
Y = np.random.randint(0, 100, (2, 2))

# Addition of matrices
result_add = X + Y
print(X)
print(Y)
```

```
print(result_add)
```

```
[[76 39]
 [76 14]]
[[65  1]
 [ 5 12]]
[[141  40]
 [ 81  26]]
```

2.3 Dividing with a Matrix

While direct division isn't defined in matrix operations, we can achieve a similar result by multiplying by the inverse of a matrix. The inverse of a matrix \mathbf{X} , denoted \mathbf{X}^{-1} , satisfies:

$$\mathbf{X} \times \mathbf{X}^{-1} = \mathbf{I} = \begin{pmatrix} 1 & 0 & \cdots & 0 \\ 0 & 1 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & 1 \end{pmatrix}$$

where \mathbf{I} is the identity matrix. Multiplying any matrix by \mathbf{I} results in the original matrix, akin to multiplying any number by 1.

In practice, while the concept is straightforward, the actual calculation of a matrix inverse can become complex for larger matrices and is typically handled by computers. We will not go through the method of obtaining the inverse in this course, we will instead just utilize the numpy function for calculating the inverse. Specifically, we use `np.linalg.inv(X)`. We can check that it actually complies with the definition like this:

```
X = np.random.randint(0, 10, (3, 3))
# Calculating inverse of X
X_inv = np.linalg.inv(X)

# Testing
np.round(np.dot(X_inv, X), 1)
```

2.4 Solving Equations with Matrix Algebra

The foundation we've established for matrix algebra now allows us to efficiently solve systems of equations. Consider solving the following pair of simultaneous equations:

$$x_{11}a_1 + x_{12}a_2 = b_1 \quad x_{21}a_1 + x_{22}a_2 = b_2$$

Here, we know the values of x and b but need to find the values of a . These equations can be succinctly expressed using matrix notation:

$$\mathbf{X} \times \mathbf{a} = \mathbf{b}$$

where \mathbf{a} and \mathbf{b} are column vectors. Let us define the right hand side vector \mathbf{b} and the coefficient matrix \mathbf{X} randomly in python as

```
b = np.random.randint(0,100,(2,1))
# Define matrix X
X = np.random.randint(0,100,(2,2))
print(X)
print(b)
```

```
[[73 79]
 [41 72]]
[[16]
 [75]]
```

To solve for \mathbf{a} , we use the inverse of \mathbf{X} , provided it exists, and multiplies it with the left and right hand sides of the equation, just as we would divide with \mathbf{X} on both sides to solve for a single equation:

$$\mathbf{X}^{-1} \times \mathbf{X} \times \mathbf{a} = \mathbf{X}^{-1}\mathbf{b}$$

Since we know that \mathbf{X}^{-1} is the solution to $\mathbf{X}^{-1} \times \mathbf{X} = \mathbf{I}$, premultiplying with \mathbf{X}^{-1} yields:

$$\mathbf{a} = \mathbf{X}^{-1}\mathbf{b}$$

Hence, we have found an easy way to solve any linear equation. We can test that it works in python. Let us first find \mathbf{a} using this approach:

```
a = np.dot(np.linalg.inv(X), b)
a
```

```
array([[ -2.36638572],
       [  2.38919187]])
```

If you get a “Singular matrix” error its because we are generating X with a few random integers, which sometimes creates unsolvable systems, so just generate X and b again.

Now we can test, if the solution for a actually works, by applying it on the original equation $X \times a = b$. This should yield the right hand side of the equation, b :

```
np.dot(X, a)
```

```
array([[16.],
       [75.]])
```

Compare this with the actual b :

```
np.dot(X, a)
```

```
array([[16.],
       [75.]])
```

Thus, we have identified an effective method to solve any system of equations, provided that X is invertible. If X cannot be inverted, it indicates that two or more equations are essentially identical, leading to an “underdetermined” system. In such cases, some equations are redundant, and there are not enough independent equations to determine the values of all variables. Remember the fundamental rule: we need an equal number of equations and unknowns to uniquely solve for each variable.

2.5 Transposing

Transposing a matrix involves swapping its rows and columns. For example, a 2×3 matrix:

$$X_{2 \times 3} = \begin{pmatrix} x_{11} & x_{12} & x_{13} \\ x_{21} & x_{22} & x_{23} \end{pmatrix}$$

transposes to:

$$\mathbf{X}'_{2 \times 3} = \begin{pmatrix} x_{11} & x_{21} \\ x_{12} & x_{22} \\ x_{13} & x_{23} \end{pmatrix}$$

where ' denotes the transposed matrix. For a column vector \mathbf{a} , transposing and then multiplying by itself, $\mathbf{a}'\mathbf{a}$, calculates the sum of squares of its components.

```
# Example of matrix transposition
X_2x3 = np.random.randint(0,100,(2,3))
X_transposed = X_2x3.T
X_transposed
```

```
array([[30, 49],
       [28, 67],
       [33, 32]])
```

Transposition is often used to conform to the requirements of matrix multiplication, where the number of columns in the first matrix must match the number of rows in the second. If this is not the case, one might transpose the first matrix to facilitate multiplication.

3 Calculus and matrices

Deriving matrices follows similar principles to deriving polynomials. For instance:

$$\frac{d(a^2\sigma^2)}{da} = 2a\sigma^2$$

applies to scalar variables, and for a matrix Σ and a column vector \mathbf{a} , we have:

$$\frac{d(\mathbf{a}'\Sigma\mathbf{a})}{d\mathbf{a}'} = 2\Sigma\mathbf{a}$$

assuming Σ is symmetric. In practical terms, the derivative with respect to \mathbf{a} here, given some values for \mathbf{a} , is

```
# Derivation with matrix and vector
a = np.random.randint(0,100,(2,1))
Sigma = np.random.randint(0,100,(2,2))
```



```
# Derivative of a' Sigma a with respect to a
derivative = 2 * np.dot(Sigma, a)
derivative

array([[ 7946],
       [14472]])
```

We can rewrite the matrix formulation in scalar form, to check that the rule is correct. The scalar form of $\mathbf{a}'\Sigma\mathbf{a}$ is

$$\mathbf{a}'\Sigma\mathbf{a} = \sum_{j=0}^N a_j \left(\sum_{i=0}^N a_i \sigma_{ij} \right)$$

You can verify that

$$\frac{d(\mathbf{a}'\Sigma\mathbf{a})}{d\mathbf{a}} = 2 \left[\sum_{i=0}^N a_i \sigma_{i0}, \dots, \sum_{i=0}^N a_i \sigma_{iN} \right]$$

4 Optimal portfolios with more than one asset

We remember from above the previous chapter that with one asset, the optimal portfolio was

$$a_{opt} = \frac{(\mu - r)}{\lambda \sigma^2}$$

From this we concluded that:

1. The more risk-averse the person is, the less they should invest.
2. The larger the expected return of the asset, the more should be invested.
3. The greater the risk associated with the asset, represented by σ^2 , the less should be invested.

Now, let us consider the optimal investments if we have more than one asset.

4.1 Optimal Portfolios with Any Number of Assets

Let us now assume that the investor in the previous section has a portfolio of N assets, not just one. Their wealth next period, assuming the entire amount is

borrowed, is then expressed in matrix notation as:

$$W_1 = \mathbf{a}'\mathbf{x} - \mathbf{1}r$$

where \mathbf{a} represents the portfolio weights, \mathbf{x} represents the returns, and $\mathbf{1}$ is a column vector of ones, such that $\mathbf{1}r$ is a column vector of the risk-free interest rate r . Recall from earlier that the investor aims to maximize the difference between expected return and variance:

$$\max_{\mathbf{a}} Z = \mathbb{E}W_1 - \lambda \frac{1}{2} \text{var}(W_1)$$

\mathbf{x} now is a column vector of many normally distributed variables with different variances and expectations. We denote the expected returns by μ_i for asset i , and the associated vector of these returns by μ . Given a portfolio \mathbf{a} , the expected return on the portfolio then becomes:

$$\mathbb{E}W_1 = \mathbf{a}'(\mathbb{E}\mathbf{x} - \mathbf{1}r) = \mathbf{a}'(\mu - \mathbf{1}r)$$

For the variance, the risk free return r is not relevant, since means are subtracted anyway. We define the covariance matrix, all the combinations of variance and covariance between the stocks as

$$\text{var } W_1 = \mathbf{a}'\Sigma\mathbf{a} = \mathbf{a}' \begin{bmatrix} \sigma_{00} & \sigma_{01} & \cdots & \sigma_{0N} \\ \sigma_{10} & \sigma_{11} & \cdots & \vdots \\ \vdots & \vdots & \ddots & \vdots \\ \sigma_{N0} & \cdots & \cdots & \sigma_{NN} \end{bmatrix} \mathbf{a}$$

Where σ_{ij} is the covariance between i and j , and σ_i^2 is the variance of asset i . This is the covariance matrix, denoted by the capital sigma, Σ .

When a vector is normally distributed we write it as $\mathbf{x} \sim N(\mu, \Sigma)$.

We have now derived expressions for $\mathbb{E}(W_1)$ and $\text{var}(W_1)$ using matrix notation. Building on the concepts from the previous lecture, we can now formulate our portfolio optimization problem as:

$$\max_{\mathbf{a}} Z = \mathbf{a}'(\mu - \mathbf{1}r) - \lambda \frac{1}{2} \mathbf{a}'\Sigma\mathbf{a}$$

Taking the derivative with respect to \mathbf{a}' yields the N first order conditions:

$$\frac{dZ}{d\mathbf{a}} = (\mu - \mathbf{1}r) - \lambda \Sigma \mathbf{a} = 0$$

Hence, in optimum:

$$\Sigma \mathbf{a} = \frac{1}{\lambda} (\mu - \mathbf{1}r)$$

By premultiplying with the inverse of Σ , we obtain the optimal portfolio:

$$\mathbf{a}_{\text{opt}} = \frac{1}{\lambda} \Sigma^{-1} (\mu - \mathbf{1}r)$$

Note that this formula looks very similar to the formula for an optimal portfolio with only one asset:

$$a_{\text{opt}} = \frac{\mu - r}{\lambda \sigma^2}$$

In general, we may draw the same conclusions as in the case of one asset:

1. The more risk-averse the person is (large λ), the less they should invest.
2. The larger the expected return the asset has, the more should be invested.
3. The more risk is associated with the asset, the less should be invested.

5 Empirical example - optimal portfolio and the portfolio front

We will create an optimal portfolio and draw the “portfolio front” - which are the smallest possible volatility of a set of assets, for all return levels. We use the script feature of Titlon to fetch the data, and storing it in 'output/stocks.df'

```
import pandas as pd
#Query script for MySQL client
import pymysql
con = pymysql.connect(host='titlon.uit.no',
                      user="user@name.com",
                      password="passwordfromtitlon",
```

```

        database='OSE')
crsr=con.cursor()
crsr.execute("SET SESSION MAX_EXECUTION_TIME=60000;")
crsr.execute("""
    SELECT  * FROM `OSE`.`equity`
    WHERE year(`Date`) >= 2016
    ORDER BY `Name`,`Date`
""")
r=crsr.fetchall()
df=pd.DataFrame(list(r),
                 columns=[i[0] for i in crsr.description])
df

#YOU NEED TO BE CONNECTED TO YOUR INSTITUTION VIA VPN,
# OR BE AT THE INSTITUTION, FOR THIS CODE TO WORK

pd.to_pickle(df, 'output/stocks.df')

```

We now reload the data, and pick the four stocks

1. has been traded at the first and last day of the sample 2. that are most traded (sorted on sum of Turnover)

We can obtain that with this code:

```

import numpy as np
import pandas as pd

df = pd.read_pickle('output/stocks.df')
# Defining annual risk free rate.
rf = df['NOWA_DayLnrate'].mean()*7

isin_with_first_date = df[df['Date'] == df['Date'].min()][['ISIN']].unique()
isin_with_last_date = df[df['Date'] == df['Date'].max()][['ISIN']].unique()
valid_isins = set(isin_with_first_date).intersection(isin_with_last_date)
df[df['ISIN'].isin(valid_isins)]

df['Name (ISIN)'] =df['Name'].str.upper() + '(' + df['ISIN'] + ')'

```

```
# keeping only the most traded shares
res = (
    df.groupby(['Name (ISIN)'])
      .agg({'Turnover': 'sum'})
      .sort_values(by='Turnover', ascending=False)
)
df = df.merge(res.head(4), on=['Name (ISIN)'],
              how='inner')
res.head(4)
```

Name (ISIN)	Turnover
EQUINOR (NO0010096985)	8.202209e+11
EQUINOR(NO0010096985)	7.030530e+11
YARA INTERNATIONAL(NO0010208051)	4.753943e+11
DNB(NO0010031479)	3.943389e+11

5.1 The historic mean and covariance matrix

We will now calculate the covariance matrix and the mean vector. To begin, we'll create a function that reformats the data from its long format (where stock prices are listed sequentially in a single column) to a wide format. In this wide format, unique dates will be in the first column, with subsequent columns containing prices for each stock.

```
def get_matrix(df, field):
    """Converts the df to a matrix df that can
    be used to calculate the covariance matrix"""

    df['Date'] = pd.to_datetime(df['Date'])
    df_unique = df.drop_duplicates(
        subset=['Date', 'ISIN'])
    pivot_df = df_unique.pivot(index='Date',
                               columns='ISIN',
                               values=field)

    pivot_df = pivot_df.dropna()
```

```
# Annualized weekly returns
df_weekly = pivot_df.resample('W').sum()

return df_weekly

#X is a matrix with e
X = get_matrix(df, 'lnDeltaP')
X.sort_index()
```

ISIN	NO0010031479	NO0010096985	NO0010208051
Date			
2016-01-10	-0.077617	-0.118288	-0.058065
2016-01-17	-0.052537	-0.060966	-0.047905
2016-01-24	0.045624	0.060966	0.001741
2016-01-31	0.032182	0.074498	-0.053584
2016-02-07	-0.017425	0.027490	0.024170
...
2020-11-01	-0.084158	-0.082292	-0.038642
2020-11-08	0.038862	0.021620	0.041243
2020-11-15	0.116952	0.087030	0.061829
2020-11-22	0.041035	0.020627	0.034672
2020-11-29	0.042786	0.068337	0.000550

With this data, it is relatively easy to calculate covariance and the means vector

```
# Converting X to a numpy array:
X = np.array(X)

# Calculating the covariance
cov_matrix = np.cov(X, rowvar=False)

# Calculating the means vector, and reshaping it to a
# column vector.

means = np.mean(X, axis=0).reshape((X.shape[1],1))
```

5.2 Plotting the portfolio front

We now turn to portfolio front. The portfolio front represents the volatility of the portfolio with the least variance, for a given portfolio return. Hence, we want a function of portfolio return that represent the minimum variance portfolios.

It turns out that by defining a few simple scalars, there is a reasonably simple expression for the set of minimum variance portfolios. The scalars are:

$$A = \mathbf{1}'\Sigma^{-1}\mathbf{1}$$

$$B = \mathbf{1}'\Sigma^{-1}\mu - \mathbf{1}r$$

$$C = \mu - \mathbf{1}r'\Sigma^{-1}\mu - \mathbf{1}r$$

With these definition, the portfolio front, representing all minimum variance portfolios for a given expected return and a total investment of 1, is

$$\sigma_{opt} = \frac{1}{A} + \frac{\left(\mu_{opt} - \frac{B}{A}\right)^2}{C - \frac{B^2}{A}}$$

where $\mu_{opt} = \mu' \mathbf{a}$

This formula automatically ensures that the sum of all the portfolio weights are 1. Now, let us define these variables, and the function. A, B and C in python:

```
ones = np.ones((len(means), 1))

A = np.dot(ones.T,
            np.dot(np.linalg.inv(cov_matrix), ones))[0][0]

B = np.dot(ones.T,
            np.dot(np.linalg.inv(cov_matrix),
                    (means-rf)))[0][0]

C = np.dot((means.T-rf),
            np.dot(np.linalg.inv(cov_matrix),
                    (means-rf)))[0][0]
```

The portfolio front function, that returns the volatility associated with the minimum variance portfolio for a given expected_excess_return can then be defined as:

```
def portfolio_front(expected_excess_return):
    r = expected_excess_return
    minimum_variance = (1/A
                        + ((r - abs(B)/A)**2) / (C - B**2/A))
    minimum_volatility = minimum_variance**0.5
    return minimum_volatility
```


Let us plot this

With a defined portfolio frontier function, we can now plot the portfolio frontier. To make the weekly returns comparable on an annual scale, we multiply them by a scaling factor, plot_scale, set to 52. We can simply multiply the returns, because returns in Titlon are conveniently defined as log returns, calculated as $\log(x_t) - \log(x_{t-1})$, rather than the typical relative return $\frac{x_t - x_{t-1}}{x_{t-1}}$.

```
from matplotlib import pyplot as plt
#Creating plot
fig, ax = plt.subplots(figsize=(10, 6))
plot_scale = 52
MAX_AXIS = 0.005
#applying the function
rp_values = np.linspace(0, MAX_AXIS-rf, 100)
sigma_values = portfolio_front(rp_values)

#plotting, after annualizing the weekly data
ax.plot(plot_scale**0.5*(sigma_values), plot_scale*(rp_values+rf),
        label='Efficient Frontier')

#plot settings:
ax.set_xlim([0, np.max(sigma_values*plot_scale**0.5)])
ax.set_ylim([0, (np.max(rp_values)+rf)*plot_scale])
ax.set_xlabel('Sigma (Risk)')
ax.set_ylabel('Rp (Return)')
ax.set_title('Efficient Frontier')
ax.legend()
```

A key advantage of using log returns is that they account for symmetry in gains and losses. For example, with relative returns, a 10% gain followed by a 10% loss results in a net loss of 1%. However, with log returns, a 10% gain and a 10% loss precisely offset each other, leaving you with no net change from the starting point.

Let us now add the point for the optimal portfolio. The optimal portfolio is

$$\mathbf{a}_{opt} = \frac{1}{\lambda} \Sigma^{-1} (\mu - \mathbf{1}r)$$

The total cost of this portfolio is

$$\mathbf{1}' \mathbf{a}_{opt} = \frac{1}{\lambda} \mathbf{1}' \Sigma^{-1} (\mu - \mathbf{1}r)$$

So the normalized portfolio is

$$\mathbf{a}_{norm} = \frac{\Sigma^{-1} (\mu - \mathbf{1}r)}{\mathbf{1}' \Sigma^{-1} (\mu - \mathbf{1}r)}$$

The expected return of the optimal portfolio on the frontier, is then

$$\mu_{port} = (\mu - \mathbf{1}r)' \mathbf{a}_{norm} = \frac{(\mu - \mathbf{1}r)' \Sigma^{-1} (\mu - \mathbf{1}r)}{\mathbf{1}' \Sigma^{-1} (\mu - \mathbf{1}r)}$$


Which is simply, according to the previous definitions of B and C:

$$\mu_{port} = \frac{C}{B}$$

Let us add that to the plot

```
# Calculating the tangency point of the normalized
# optimal portfolio
tangency_sigma = portfolio_front(C/B)

#plotting it, after annualizing the weekly data
ax.plot(plot_scale**0.5*tangency_sigma,
        plot_scale*(C/B + rf),
        'ro',label='Tangency Point')
ax.legend()
fig
```



Let us now draw a tangency line from the risk free interest rate `rf` to the optimal portfolio point. The slope will be the Sharp-ratio

$$S = \frac{C/B}{f(C/B)}$$

where $f(C/B)$ is the portfolio front function `portfolio_front(C/B)` at C/B


```
import pickle #For saving the plot

sigma_range = np.linspace(0, np.max(sigma_values), 100)

# Plotting the portfolio front, after annualizing the
# weekly data
ax.plot(plot_scale**0.5*sigma_range, plot_scale*(rf
        + sigma_range*(C/B)/tangency_sigma),
        color='r', linestyle='--',
        label='Capital Market Line')

ax.legend()

fig
```



Coding Challenges:

- **Challenge 1:** Obtain a three or four return series for portfolios, for example your portfolio, the market portfolio and the factors at titlon and calculate the variance-covariance matrix and the means
- **Challenge 2:** Use your calculated covariance matrix and means, and the code provided above, to draw a portfolio front.
- **Challenge 3:** Use the same information to plot the points of each portfolio in the same chart.
- **Challenge 4:** Calculate the optimal portfolio and place it in the chart, together with the capital market line

6 Factors

Factors are portfolios constructed based on specific characteristics of assets. For example, Fama and French proposed factors based on company valuation (over- or under-valuation) and company size.

6.1 Construction of Factors

A standard, simplified method for constructing a factor portfolio involves creating a long position in the third of assets with the strongest characteristic (e.g., the most undervalued) and a short position in the third with the weakest. For instance, the **HML** (High Minus Low) factor by Fama and French takes a long position in the top third of companies with the highest market-to-book value and shorts the bottom third with the lowest.

This results in a portfolio with zero net cost because the long and short positions offset each other. Hence, there's no need to subtract the risk-free rate when calculating returns for factor portfolios.

In Titlon, there are four factors: **SMB**, **HML**, **LIQ**, and **MOM**, alongside the market factor (the market index).

- **SMB (Small Minus Big):** The return of small companies minus the return of large companies.
- **HML (High Minus Low):** The return of companies with high market-to-book ratios compared to those with low ratios.
- **LIQ (Liquidity):** The return of the most liquid companies minus that of the least liquid.

- **MOM (Momentum):** The return of companies with high momentum minus those with high reversal tendencies.

6.2 Factor Model

The main purpose of factors is to be used in regression analysis, like the following:

$$\begin{aligned} x - r_f = & \alpha + \beta_M(r_M - r_f) \\ & + \beta_{SMB} \cdot SMB \\ & + \beta_{HML} \cdot HML + \beta_{LIQ} \cdot LIQ \\ & + \beta_{MOM} \cdot MOM + \epsilon \end{aligned}$$

This is a multifactor model. If we only include the market factor, the model reduces to the well-known **CAPM** (Capital Asset Pricing Model). Adding the additional factors results in a more comprehensive **factor model**.

6.3 Factors in Portfolio Evaluation

Numerous factors have been proposed in the literature. However, many are believed to be the result of data mining, so it's common practice to use only the most established ones, like those mentioned above, when evaluating portfolio performance.

The estimated α from the factor model is the most widely recognized measure of risk-adjusted return. A positive α indicates that a portion of the excess return of $x - r_f$ cannot be explained by exposure to any of the factors, including the market factor. This implies that the portfolio has delivered some form of risk-free excess return.

As in any regression, you can compute the standard error and p-value of the estimated α . This is crucial because, if the α is not statistically significant, we cannot confidently conclude that it is different from zero. Therefore, to claim that an asset or portfolio has truly outperformed the market, its multifactor α should be both positive and statistically significant.

6.4 Historical Context

This framework is the standard method for determining whether a portfolio manager has genuinely been skilled or simply benefited from luck or factor exposures.

The field of finance has, in many ways, been driven by the need to explain portfolio managers' overperformance. In the early 20th century, some managers appeared to consistently outperform the market. The development of **CAPM** revealed that this was often due to selecting stocks with high market risk rather than genuine skill. Most of these managers did not generate CAPM alpha.

For investors seeking higher returns by taking on more market risk, simply buying more shares (increasing exposure to the market) is typically more cost-effective than picking the riskiest stocks. The development of CAPM helped investors make more informed choices.

Similarly, the introduction of factor models showed that some managers were merely betting on small-cap stocks or undervalued companies to generate excess returns. When accounting for these factors, much of the supposed excess return often disappears.

6.4.1 Coding Challenges:

- **Challenge 1:**
 - Download stock data from Titlon for a single stock.
 - Perform a multifactor regression model using the downloaded data.
 - Analyze the significance of the alpha: check whether it is statistically significant, and interpret its direction (positive or negative).
 - Provide commentary on what the result implies in terms of the stock's performance relative to the factors.
- **Challenge 2:**
 - Download factor data from Titlon and plot the performance of these factors alongside the optimal portfolio from the previous chapter.
 - Add each factor as a data point on the chart, allowing comparison between the factors and the optimal portfolio as well as the portfolio frontier.
 - Use the following SQL query in the Titlon script to retrieve the factor data:

```
SELECT [SMB], [HML], [LIQ], [MOM]
FROM [OSE].[dbo].[factors]
WHERE YEAR([Date]) >= 2016
```

- Plot the performance and visually assess how each factor performs relative to the optimal portfolio.

6.4.2 Tentative answer to Challenge 2:

We first download the factor data from titlon

```
import pandas as pd
#Query script for MySQL client
import pymysql
con = pymysql.connect(host='titlon.uit.no',
                      user      = "user@account.com",
                      password   = "%VccvGWz5UjtT9sAuU7og",
                      database='OSE')
crsr=con.cursor()
crsr.execute("SET SESSION MAX_EXECUTION_TIME=60000;")
crsr.execute("""
    SELECT  *
    FROM `OSE`.`factors`
    WHERE year(`Date`) >= 2016
    ORDER BY `Date`
""")
r=crsr.fetchall()
df=pd.DataFrame(list(r),
                columns=[i[0] for i in crsr.description])
df

#YOU NEED TO BE CONNECTED TO YOUR INSTITUTION VIA VPN,
# OR BE AT THE INSTITUTION, FOR THIS CODE TO WORK

pd.to_pickle(df, 'output/factors.df')
```

We then calculated the means and standard errors, after reducing the frequency from daily to weekly.

```
import pandas as pd

df = pd.read_pickle('output/factors.df')
df['Date'] = pd.to_datetime(df['Date'])
df = df.set_index('Date')
df_weekly = df.resample('W').sum()
df_weekly = df_weekly[['SMB', 'HML', 'LIQ', 'MOM']].dropna()
df = df[['SMB', 'HML', 'LIQ', 'MOM']].dropna()
```

```

means = df_weekly.mean()*52
std = df_weekly.std()*52**0.5
print(std)
print(means)
df_weekly

```

```

SMB      0.145195
HML      0.193652
LIQ      0.160615
MOM      0.216708
dtype: float64
SMB      0.074676
HML      0.145056
LIQ      0.019353
MOM      0.164041
dtype: float64

```

	SMB	HML	LIQ	MOM
Date				
2016-01-10	0.040139	-0.038205	0.063480	0.039485
2016-01-17	-0.004794	-0.053537	0.045607	0.013392
2016-01-24	0.016701	0.025072	-0.004320	-0.014409
2016-01-31	0.002747	-0.001928	-0.016214	-0.024070
2016-02-07	-0.008014	-0.029920	-0.002909	-0.008874
...
2024-03-10	0.032518	0.037676	0.017710	-0.005075
2024-03-17	0.014078	0.007018	-0.025601	0.023273
2024-03-24	0.013395	-0.008679	-0.003795	0.036296
2024-03-31	0.000000	0.000000	0.000000	0.000000
2024-04-07	0.000000	0.000000	0.000000	0.000000

We can then plot the factors in the existing portfolio frontier


```

for series in means.index:
    ax.scatter(std[series], means[series], label=series)

ax.legend()

```


fig



Can you construct the optimal portfolio of the stocks and the factors?